

# Introduction to OpenMP

- Goal: make parallel programming easy
- Shared memory programming
- Compiler pragma for C/C++ and Fortran
  - Requires a compatible compiler
- Require a runtime helper library
- Portable solution
- Official site: <http://openmp.org/>

# Parallel section example

#pragma omp: compiler directive

```
/*  
 * Example 1: omp parallel  
 */  
QDebug() << "omp parallel begin";  
#pragma omp parallel  
{  
    QDebug() << "rank" << omp_get_thread_num() <<  
        "size" << omp_get_num_threads();  
}  
QDebug() << "omp parallel end";
```

API OpenMP

There is an implicit barrier at the end of a parallel scope

Chaque fil d'exécution va exécuter le bloc de code.

# pragma

- `#pragma omp directive [clause, ...]`
- Applies to a scope `{ ... }`
  - or the following statement
- Directives:
  - parallel, critical, single, barrier, atomic
- Clauses:
  - for, private, reduction

# API OpenMP

- Utilities (optional)
- `#include <omp.h>`
- Examples:
  - `omp_get_num_threads()`: returns the number of started threads
  - `omp_get_thread_num()` returns the thread rank

# Environnement variables

- Can modify the behavior without recompiling
- Example:
  - OMP\_WAIT\_POLICY: passive or active wait on a barrier
  - OMP\_NUM\_THREADS: number of threads to spawn (default is the number of cores)

# Compiling OpenMP program

- GCC: -fopenmp (man gcc)
- Compile: QMAKE\_CXXFLAGS += -fopenmp
- Link: QMAKE\_LFLAGS += -fopenmp
- warning: ignoring #pragma omp parallel [-Wunknown-pragmas]
- error: undefined reference to `omp\_get\_num\_threads'

# Parallel computation

```
auto memset_serial = [&]() {  
    for (int i = 0; i < n; i++) {  
        v[i] = i * i;  
    }  
};
```

```
auto memset_parallel = [&]() {  
    #pragma omp parallel for  
    for (int i = 0; i < n; i++) {  
        v[i] = i * i;  
    }  
};
```

#pragma omp parallel:  
spawns a parallel section

#pragma omp for:  
Split loop iterations

# Private or shared variables?

```
// variables
int n = 1000;
float max = 0.0;
int i = 0;
float p = 0.1;
float t = 0.0;
QVector<float> orig(n);
QVector<float> norm(n);
```

Parent scope: shared by default

One exception: loop index is private by default

```
// init sound-like data
#pragma omp parallel for
for (i = 0; i < n; i++) {
    orig[i] = std::cos((t + 0.1) * 0.1) +
              std::cos((t + 0.2) * 0.2) +
              std::cos((t + 0.3) * 0.3);

    t += p;
}
```

Beware: t is shared and subject to race condition. Must protect it with mutex (or #pragma omp critical) or use private(t) or some reduction operator.



# Variable scope

- Clause shared: shared by default, except loop index
- Clause private: thread local
- Clause firstprivate: thread local, but initialized with the content of the parent scope
- Also available: lastprivate, copyprivate

# Critical section

```
// find the max value (manual reduction)
float max = 0.0;
float my_max = 0.0;
#pragma omp parallel private(my_max)
{
    #pragma omp for
    for (i = 0; i < n; i++) {
        my_max = std::max(my_max, std::abs(norm[i]));
    }
    #pragma omp critical
    {
        max = std::max(max, my_max);
    }
}

QDebug() << "check max=" << max;
```

Variable my\_max private to each thread

Safe update of shared variable.

# Reduction

Variable the\_max shared,  
but OpenMP creates a local  
variable for intermediate reduction

```
// find the max value (automatic reduction)
float the_max = 0.0;
#pragma omp parallel for reduction(max:the_max)
for (i = 0; i < n; i++) {
    the_max = std::max(the_max, std::abs(norm[i]));
}
```

```
qDebug() << "check max=" << the_max;
```

Implicit critical section for the final  
reduction.

Operators for reduction (initialization values)			
+	(0)		(0)
*	(1)	^	(0)
-	(0)	&&	(1)
&	(~0)		(0)
max (Least representable number in <b>reduction</b> list item type)			
min (Largest representable number in <b>reduction</b> list item type)			

Source: OpenMP 4.0 API C/C++ Syntax Quick Reference Card

# Independent computation: nowait

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (int i=0; i<n; i++) a[i] = b[i] + c[i];

    #pragma omp for nowait
    for (int i=0; i<n; i++) d[i] = e[i] + f[i];

    #pragma omp barrier
    scale = sum(a,0,n) + sum(d,0,n);
}
```

Explicit barrier required to wait for all previous computation to finish

The first and the second loops execute at the same time

# Examples

- 19-omp-base: pragma examples
- 20-omp-speedup: Measure speedup obtained with OpenMP
- 21-omp-scope: Examples of variable scope
- 22-omp-seq: Using OpenMP to normalize sound signal
- 23-omp-task: Traverse tree-like structure in parallel
- 24-omp-atomic: Example of pragma atomic

# Performance counters


- Counts microarchitecture events
- Cache accesses and misses
  - load, store
  - L1, Last Level Cache (LLC)
  - data, instruction
- Can compute cache misses ratio
- Can compute  $\text{Cycle} / \text{Instructions} = \text{CPI}$

# Cache misses

- Linear access cold start
  - Few first accesses causes cache misses, then hardware pre-fetcher will populate the cache in advance
- Large stride in memory (accessing 2D data by column)
- Random access
- Data size v.s. size of the cache: *working set size*
- Sharing a cache line with other cores
  - False sharing and cache line bouncing (lock)

# False sharing

```
typedef struct point {  
    int x;  
    int y;  
} point_t;
```



```
point_t *foo;
```

## Modification (CPU 0)

```
// moving in an airplane  
// wheeee!
```

```
for (;;)   
    foo->x++;
```

## Read (CPU 1)

```
// if latitude is close to the equator  
// then you are lucky
```

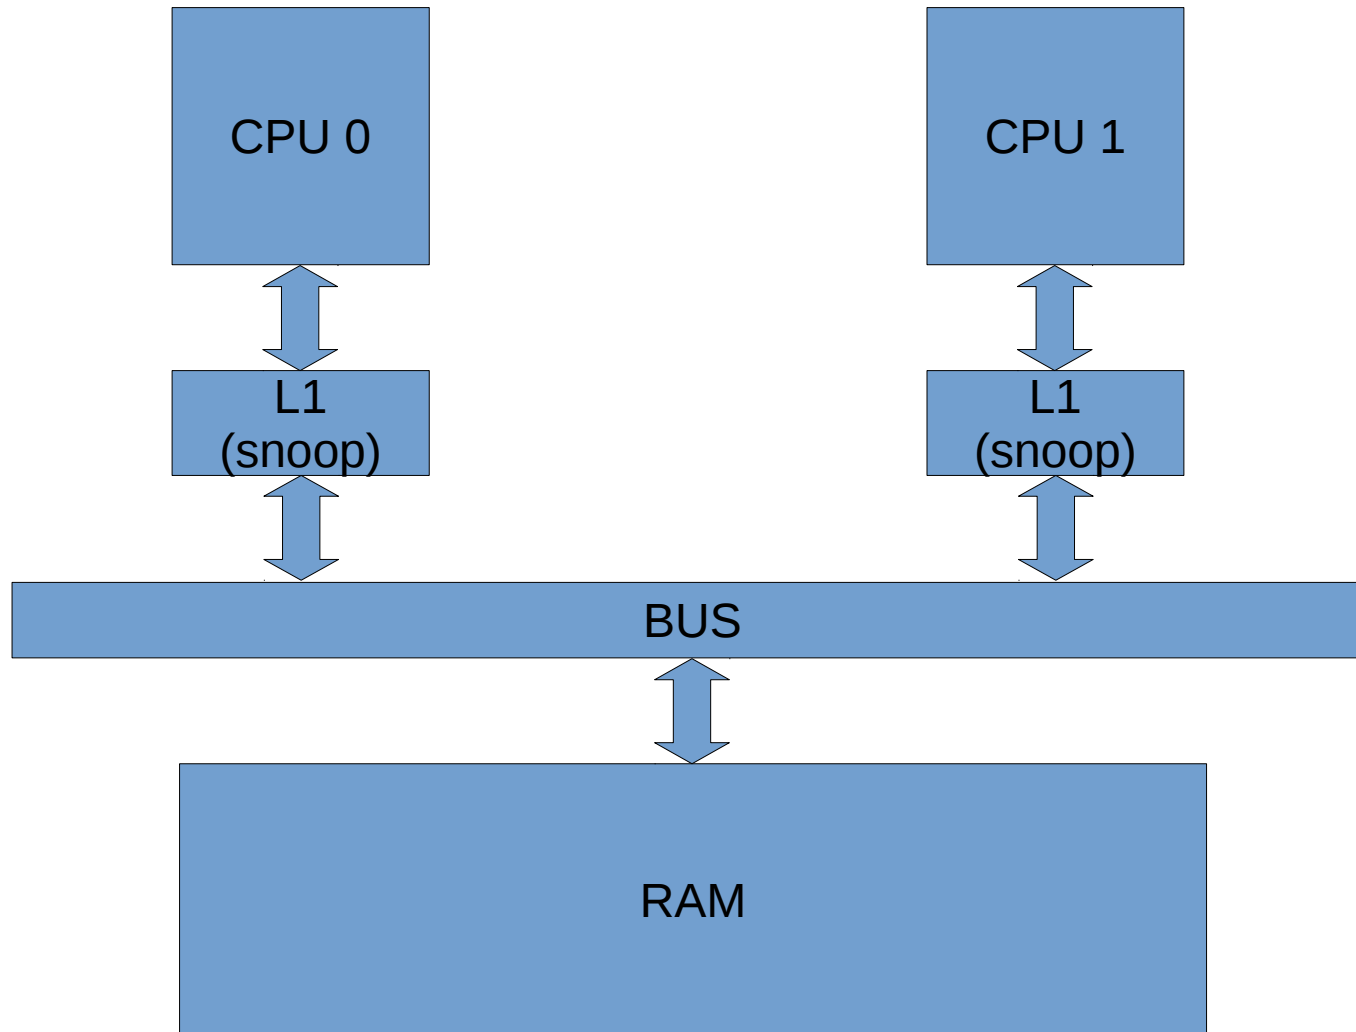
```
bool is_lucky(point_t &p) {  
    return (foo->y > -20 &&  
            foo->y < -20);  
}
```

```
for (;;)   
    is_lucky(foo);
```



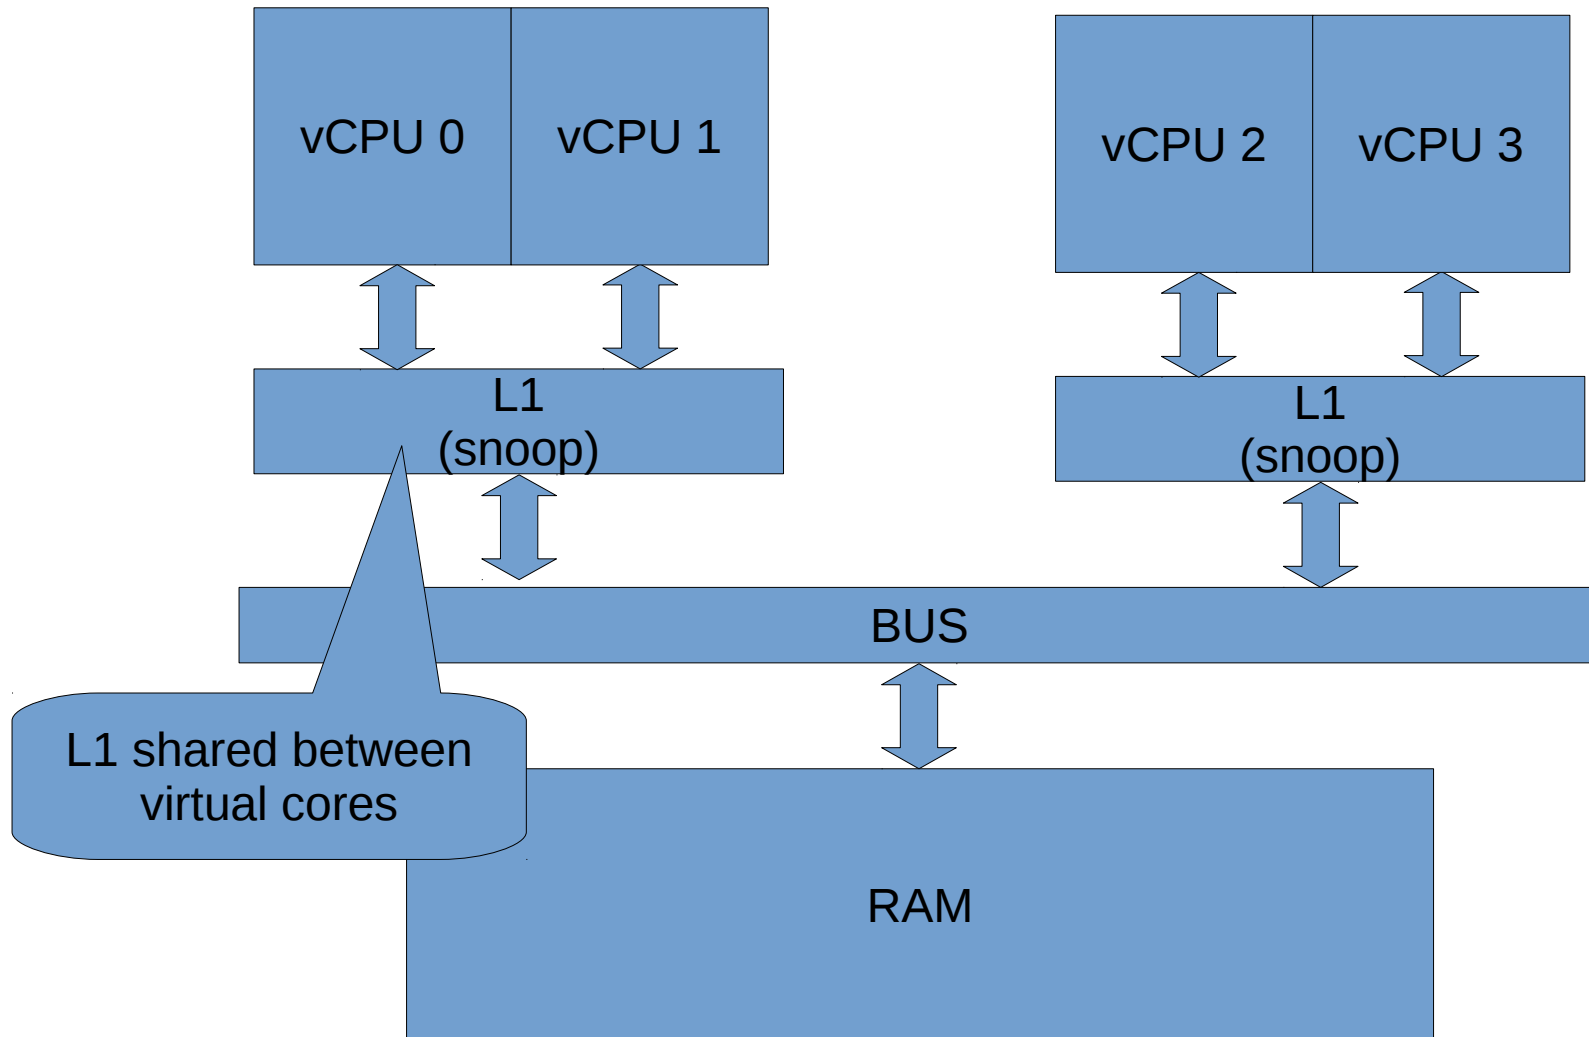
# Simultaneous Multithreading (SMT)

Intel: HyperThreading



# Simultaneous Multithreading (SMT)

Intel: HyperThreading



# Activity encode

- Encode applies rot-13 on a buffer and compute the checksum
- Implement the `encode_fast()` with OpenMP
- Run benchmarks:
  - `./perf-encode.sh pmu`
  - `./perf-encode time`
- Load the data in the spreadsheet `encode.ods`
  - Each result file should be copied into the corresponding sheet based on its name
- Analyze the data